

PeriDefT: a periodic-Fourier Transformer for 2D defect formation energy, validated by prospective first-principles DFT

Anonymous

(authors and affiliations to be filled in for submission)

Draft v3.1 May 7, 2026

Abstract

Most published machine-learning surrogates for 2D defect formation energy E_f are benchmarked exclusively on a held-out test fold of the same database used for training. This evaluation overstates deployment utility because the chemistry of interest in any materials-discovery loop is by construction out of distribution. We address this gap with three contributions. *First*, we introduce PERIDEFT, a compact 0.75 M-parameter hybrid GNN-Transformer that augments self-attention with a translation-invariant **Periodic Fourier Bias** (PFA) on minimum-image fractional displacements and trains jointly across four DFT databases (IMP2D + JARVIS-2D + JARVIS-3D + JARVIS DFT-3D, 26 837 samples) through per-source readout heads. PFA’s exact lattice-translation invariance is what enables stable joint training across geometrically heterogeneous sources — a coupling we verify with a matched single-source ablation. PERIDEFT achieves a 4-seed IMP2D test MAE of **0.486 ± 0.025 eV**, competitive with ALIGNN (4.03 M parameters, 0.540 eV) at 1/5 the parameter budget; a 6-seed temperature-calibrated ensemble reaches **0.443 eV** with 90% interval coverage of 93.4%. *Second*, we run **60 first-principles PBE-DFT calculations on the model’s own out-of-distribution predictions**, split into a “low- E_f confident” bucket (top-30 by predicted μ) and a “high- σ stress-test” bucket (top-30 by calibrated uncertainty). On the 37 SCF-converged candidates, the overall predicted-DFT Pearson correlation is +0.354 (95 % bootstrap CI [+0.06, +0.70], excludes zero), and **14/20 = 70 % of bucket-A picks are DFT-confirmed at $E_f < +1$ eV** (95 % CI [50%, 90%]), $\sim 4\times$ the predicted-pool base rate. Within the OOD bucket B (n=17), the within-bucket Pearson(σ_{cal} , |error|) is -0.288 (95 % CI $[-0.60, +0.09]$): the central tendency is mildly negative but the CI weakly straddles zero, placing a candidate deployment caveat that σ_{cal} may be useful as an in-vs-OOD bucket gate but not as a sample-level reliability score on truly unseen chemistry. On this 37-candidate OOD subset specifically, the per-dopant-corrected MAE is $\sim 5.5\times$ the in-distribution IMP2D test MAE, a concrete deployment-relevant gap that we do not extrapolate to a universal bound. *Third*, we release the full pipeline, including 125 Quantum-ESPRESSO SCF outputs, leak-free augmentation contracts, and the calibrated ensemble, to enable direct reuse for downstream 2D defect screening.

Keywords: 2D materials, defect formation energy, graph neural networks, self-attention, periodic Fourier bias, multi-source training, calibrated uncertainty, prospective DFT validation.

1 Introduction

Two-dimensional (2D) materials owe their distinctive electronic, optical and catalytic properties to atomic-scale point defects — vacancies, interstitials, substitutionals, adatoms, and adsorbed species. Defect engineering is now a central paradigm in 2D materials design [7, 12], and the

central thermodynamic quantity that controls which defects are present in equilibrium is the *formation energy*

$$E_f(D, q) = E_{\text{tot}}(D, q) - E_{\text{tot}}(\text{host}) - \sum_i n_i \mu_i + q(E_F + E_{\text{VBM}}) + E_{\text{corr}}, \quad (1)$$

where $E_{\text{tot}}(D, q)$ is the total density-functional-theory (DFT) energy of the defect supercell at charge state q , $E_{\text{tot}}(\text{host})$ is the host pristine-supercell energy, μ_i are the chemical potentials of species added or removed, and E_{corr} collects finite-size and image-charge corrections. Computing E_f from first principles requires fully relaxed defect supercells with careful chemical-potential bookkeeping; the cost (tens to hundreds of CPU-hours per configuration) makes high-throughput exploration of the (host, dopant, site) design space infeasible without machine-learning surrogates.

Existing graph neural networks (CGCNN [17], SchNet [13], ALIGNN [5], M3GNet [4]) and Transformer extensions (Matformer [18], Crystalformer [15], PotNet [9]) have driven 2D pristine-material property prediction to within ~ 0.05 eV/atom of DFT. For point defects in 2D, however, two complications appear. *First*, the perturbation introduced by a defect is local in space but non-local in its electronic and elastic response, so models with strict short-range cut-offs underestimate the response field [2]. *Second*, the available training data per defect type are sparse compared to bulk crystal databases; in our principal source (IMP2D, $\sim 10^4$ samples) a previously reported empirical scaling law indicates [1]

$$\log \text{MAE} \approx 3.39 - 0.40 \log N_{\text{train}} - 0.01 \log N_{\text{params}}, \quad R^2 = 0.95, \quad (2)$$

i.e. test accuracy is *very* sensitive to training-set size and *nearly insensitive* to model capacity in the regime we operate. This poses a sharp methodological question: *what is the relative payoff between architecture engineering and data engineering at the present data scale?*

A second — and more consequential — gap lies in how 2D-defect ML models are evaluated. Almost every published surrogate for E_f reports a held-out test-fold MAE on the same database used for training. Yet the chemistry of practical interest in any materials-discovery application (e.g. a new dopant on a known host, or a new host on a known dopant) is by construction *out of distribution*, and a test-fold MAE provides almost no information about the surrogate’s deployment utility. Closing this gap requires what happens routinely in biology and protein modelling but rarely in 2D-defect ML: *prospective* validation against fresh first-principles calculations on the model’s own picks.

Contributions. This paper makes three contributions along these two gaps.

1. We introduce **PERIDEFT**, a compact 0.75 M-parameter hybrid GNN–Transformer for 2D defect formation energy. The encoder augments standard self-attention with a *Periodic Fourier Bias* (PFA) on minimum-image fractional displacements, making the attention scores exactly invariant under lattice translations of any single atom. PFA’s translation invariance is what enables stable joint training across geometrically heterogeneous DFT databases via per-source readout heads (IMP2D + JARVIS-2D + JARVIS-3D + JARVIS DFT-3D, 26 837 samples, $3.15 \times \text{IMP2D}$), a coupling we verify with a matched single-source ablation. The 4-seed multi-source PERIDEFT achieves IMP2D test MAE **0.486 $\pm$ 0.025** eV, an 11 % improvement over a matched-architecture single-source baseline and competitive with ALIGNN (4.03 M parameters, 0.540 eV) at 1/5 the parameter budget; a 6-seed temperature-scaled ensemble reaches **0.443** eV with 90 % interval coverage of 93.4 %.
2. We report the first **prospective DFT validation** of an IMP2D-trained 2D-defect E_f model with calibrated uncertainty: 60 first-principles PBE-DFT calculations on the model’s own out-of-distribution predictions, split into a “low- E_f confident” bucket (top-30 by predicted

μ , “discovery”) and a “high- σ stress-test” bucket (top-30 by calibrated σ_{cal}). On the 37 SCF-converged candidates, **14/20 = 70 % of bucket-A picks are DFT-confirmed at $E_f < +1$ eV** (95 % CI [50%, 90%]), $\sim 4\times$ the predicted-pool base rate, with the overall predicted–DFT Pearson correlation +0.354 (95 % CI [+0.06, +0.70], excludes zero). Within the OOD bucket B, $\text{Pearson}(\sigma_{\text{cal}}, |\text{error}|) = -0.288$ (95 % CI [−0.60, +0.09]) — the point estimate suggests σ_{cal} may be informative for in-vs-OOD bucket triage but not for sample-level error ranking on OOD, though the within-bucket sample size warrants confirmation at larger N . On this 37-candidate OOD subset, the per-dopant-corrected MAE is $\sim 5.5\times$ the in-distribution test MAE: a concrete deployment-relevant gap, not a universal upper bound.

3. We release the full pipeline as **open infrastructure**: the trained PERIDEFT ensemble, leak-free augmentation contracts (including a self-caught augmentation leak audit, Sec. 5.6), the 287-candidate generation script, the 60-candidate selection record, and **125 Quantum-ESPRESSO SCF outputs** (38 atomic chemical-potential references, 27 pristine hosts, 60 defect supercells), enabling direct reuse for downstream 2D defect screening.

We further situate PERIDEFT in two retrospective controls. A matched single-source architectural ablation (Sec. 5.2) demonstrates that PFA, multi-scale distance bias and a dual-stream pristine–defect cross-attention are all statistically indistinguishable from a tuned baseline at fixed $N_{\text{train}} = 8512$ — consistent with our previously reported empirical scaling law (Eq. 2) and the prediction $\partial\text{MAE}/\partial N_{\text{train}} \gg \partial\text{MAE}/\partial(\text{arch.})$ in the single-source regime. The 11 % gain delivered by the multi-source PERIDEFT recipe is therefore not a 9th attempt at architecture engineering, but the first time the architectural lever is given the data scale to express itself. We also conduct a partial leave-one-host-out evaluation (Sec. 5.7) and four quantitative physical-interpretability tests (Sec. 5.5) to characterise when, and why, PERIDEFT generalises.

2 Related work

GNNs for crystal property prediction. CGCNN [17] represents a crystal as a graph with fixed-radius edges; SchNet [13] introduces continuous-filter convolutions; ALIGNN [5] jointly models atom and bond-line graphs for angle awareness; M3GNet [4] generalises to universal interatomic potentials. These methods all face the same challenge for defect supercells: the fixed real-space cut-off truncates the long-range elastic / electronic field that a defect introduces.

Transformers for crystals, and what PERIDEFT adds. Matformer [18] encodes periodicity through explicit periodic-image nodes; Crystalformer [15] introduces an infinitely-connected attention via reciprocal-space expansion; PotNet [9] uses an Ewald-style summation to approximate the full inter-atomic interaction. PERIDEFT’s *Periodic Fourier Bias* (PFA) is closest in spirit to Crystalformer / PotNet — both expand per-pair interaction in a reciprocal-space basis — but differs in three ways that matter for our setting.

(1) Where Crystalformer and PotNet *replace* the dense self-attention, PFA enters as an *additive bias* on vanilla scaled-dot-product attention. The ablation cost (turning PFA off) is a single-line code change, the parameter overhead is ~ 2 k, and the rest of the architecture (layer norm, residual, feed-forward) is the same as a standard Crystal Transformer. This makes the controlled comparison of Sec. 5.2 possible.

(2) PFA’s translation invariance is *exact* by the 2π -periodicity of $\cos/\sin$ on the minimum-image fractional displacement, not approximate via finite truncation. This matters specifically for the multi-source training recipe of Sec. 5.4: the four databases (IMP2D + JARVIS-2D + JARVIS-3D + JARVIS DFT-3D) use different lattice conventions and different cell sizes; an approximately periodic encoder leaks cell-convention artefacts into the readout, but PFA does not. This coupling — exact translation invariance enabling stable multi-source training through per-source readout heads — is what we identify as the architecturally distinctive design choice of

PERIDEFT, distinct from both “yet another Crystal Transformer” and a generic transfer-learning recipe.

(3) PERIDEFT is, to our knowledge, the first 2D-defect E_f surrogate that bundles a Crystal-Transformer encoder, a 6-seed temperature-scaled calibrated σ_{cal} , and a prospective first-principles DFT validation on the model’s own out-of-distribution predictions. Crystalformer / Matformer / PotNet each report test-fold MAE without uncertainty calibration, and none has been validated against fresh DFT on chemistry outside the training distribution. The contribution of PERIDEFT is therefore not a single architectural novelty but a coupled encoder–readout–validation *recipe* for 2D defect-property deployment.

Δ -learning and dual-stream architectures. Predicting a property as a difference from a cheap reference state is the Δ -*learning* idea: ANI-1ccx [14] predicts CCSD(T) total energies as small corrections to a DFT reference; Bogojeski et al. [3] apply Δ -learning to molecular kinetics. Our dual-stream architecture is a defect-physics adaptation: the “cheap reference” is the host pristine supercell, and the “correction” is conditioned on the local perturbation through cross-attention.

Scaling laws for materials ML. Empirical scaling laws have been recently reported for foundation models on chemistry datasets [11]. Our v1.2 report [1] fits a two-parameter scaling law (Eq. 2) on a 5×4 grid of $(N_{\text{train}}, N_{\text{params}})$ for the present architecture; the present paper provides the first *prospective* test of that law on an unseen N_{train} regime (multi-source training).

Defect-aware GNN architectures. Wu et al. [16] propose a graph Transformer for defect projected DOS; Hua and Lin [8] build a local-global associative frame for symmetry preservation. Our prior work [1] explored a virtual-defect-token architecture (“DAST”) and reported that it degrades performance by 0.6–1.0 eV on IMP2D. The dual-stream design we test here is a more conservative defect-aware idea: it keeps every node a real atom and only adds a cross-attention pathway, so it can be ablated cleanly.

3 Datasets and methodology

3.1 Datasets

The principal training and test set is the *Impurities in 2D Materials Database* (IMP2D) [10], comprising 17 364 2D-supercell defect configurations from PBE calculations. Following the cleaning rules of our prior work [1] we keep `converged=True` entries with $|E_f| \leq 20$ eV, giving 10 641 valid configurations spanning 44 host materials and 65 dopant species. Defect types are restricted to adsorbates (65%) and interstitials (35%); no charged defects or substitutionals. For multi-source training we add three external sources from JARVIS-DFT [6]: 70 2D vacancy structures (JARVIS-2D), 381 3D bulk vacancies (JARVIS-3D), and 19 902 3D pristine bulk formation energies (JARVIS DFT-3D). The combined corpus is 26 837 configurations across four reference conventions.

3.2 Leak-free data augmentation

For every IMP2D training sample we generate two physics-preserving augmentations: (i) a uniform random in-plane rotation $R \in \text{SO}(2)$ applied to both atomic positions and the lattice basis (rotation invariance of E_f is exact); and (ii) Gaussian coordinate perturbation with $\sigma = 0.02$ Å (a thermal-jitter analogue with negligible effect on E_f). The augmentations are applied *only after* the train/val/test split is fixed.

The aug-then-split trap. The naive “concatenate-then-split” protocol — common in the materials-ML literature — inadvertently leaks any single original sample into all three folds and inflates apparent test accuracy by a factor of $2.5\times$ on this dataset (test MAE 0.516 *vs.* 0.206 for the same checkpoint, depending on the split protocol). Re-evaluating the same checkpoint on the leak-free vs. leaked test sets quantifies this. We therefore enforce *split-then-augment* everywhere and encode the convention in a metadata flag (`n_train`, `n_val`, `n_test`) that the split utility checks. Section 5.6 describes the leak path more carefully and documents how, during the present work, we briefly re-introduced the bug ourselves under a refactor; this self-audit illustrates how easy the trap is and provides a defensible test for any downstream ML benchmark.

3.3 Atomic features and graph construction

Each atom i is represented by a 9-dimensional descriptor (group, period, Pauling electronegativity, covalent and van der Waals radii, valence-electron count, first ionisation energy, electron affinity, atomic mass), column-wise min-max normalised. For every supercell we build a sparse local graph with cut-off $5\text{ \AA}$ (real-space neighbour list with periodic-image shifts), accumulating directed edges (i, j) together with their fractional shifts and Euclidean distances. Triplets (j, i, k) are sub-sampled (32 per centre) and exposed as bond angles. We additionally retain the dense $N \times N$ minimum-image distance matrix and the lattice $\mathbf{C} \in \mathbb{R}^{3 \times 3}$ for use by the global self-attention block.

4 Architectures

4.1 v1 baseline: CrystalTransformer

The v1 baseline of [1] stacks three SchNet-style local-interaction layers (continuous-filter convolution + RBF-encoded angle modulation) followed by two global Transformer blocks. Each global block is a vanilla pre-norm self-attention with multi-head Q/K/V projections and a learnable distance bias $b_{ij}^{(h)} = \text{MLP}_h(\text{RBF}(d_{ij}))$ on the minimum-image scalar distance d_{ij} . With hidden dimension 128, four heads, and dropout 0.1, the v1 baseline has 0.747M parameters.

4.2 Periodic Fourier Bias (PFA)

Given padded atomic positions $\mathbf{r} \in \mathbb{R}^{B \times N \times 3}$ and per-sample lattice $\mathbf{C} \in \mathbb{R}^{B \times 3 \times 3}$, we compute the minimum-image fractional displacement

$$\mathbf{f}_{ij} = ((\mathbf{r}_i - \mathbf{r}_j) \mathbf{C}^{-1}) - \text{round}((\mathbf{r}_i - \mathbf{r}_j) \mathbf{C}^{-1}), \quad \mathbf{f}_{ij} \in [-\frac{1}{2}, \frac{1}{2})^3. \quad (3)$$

We then evaluate, for every pair (i, j) , a truncated Fourier basis

$$\{1, \cos(2\pi \mathbf{k} \cdot \mathbf{f}_{ij}), \sin(2\pi \mathbf{k} \cdot \mathbf{f}_{ij})\}_{\mathbf{k} \in \mathcal{K}}, \quad (4)$$

where $\mathcal{K} \subset \mathbb{Z}^3$ contains all non-zero integer triples with $\|\mathbf{k}\|^2 \leq 4$, taken in canonical (positive lexicographic) representatives to remove the $\mathbf{k} \leftrightarrow -\mathbf{k}$ redundancy of $\cos$. A small two-layer MLP projects this basis to per-head additive bias scalars $b_{ij}^{(h)}$.

Invariance. Shifting any single atom $\mathbf{r}_i$ by an integer combination of lattice vectors $\mathbf{L} \in \mathbb{Z}^3 \cdot \mathbf{C}$ leaves $\mathbf{f}_{ij}$ invariant modulo 1, and since $\cos$ and $\sin$ are 2π -periodic, $b_{ij}^{(h)}$ is exactly unchanged. The bias is also rotation- and translation-invariant by the same algebra. We provide an automated unit test (`scripts/test_attention_v2.py`) that empirically verifies all three invariances to floating-point precision (max deviation $< 10^{-7}$ for Eq. 3, end-to-end model-output deviation $< 10^{-3}$ under arbitrary rotations).

4.3 Multi-scale distance bias

The single-RBF distance bias of the v1 baseline is replaced with two channels: (i) a short-range Gaussian RBF (32 centres, $r \in [0, 5]$ Å) projected to per-head bias and tapered with a smooth cut-off near $r = 5$ Å; and (ii) a long-range analytical channel built on $[1/(r + \delta), 1/(r + \delta)^2, 1/(r + \delta)^3, e^{-r/\lambda_k}]_k$ with $\delta = 1$ Å and four learnable $\lambda_k \in [2, 10]$ Å, smoothly tapered to zero at $r_{\max} = 12$ Å. The two channels are summed.

4.4 Pristine–defect dual-stream cross-attention

The formation energy (Eq. 1) is, by physical definition, a *difference* between a defect-supercell energy and a host-pristine-supercell energy. Predicting E_f from the defect supercell alone forces the network to implicitly recover $E_{\text{tot}}(\text{host})$ from a single structure, which wastes capacity on a quantity that depends only on the chemistry and not the defect. We instead make the host pristine reference an explicit input: for every IMP2D sample we construct an *unrelaxed* pristine supercell by removing the dopant atom from the defect supercell. Because IMP2D defects are exclusively adsorbate (65%) or interstitial (35%), this construction is exact in chemistry and unrelaxed in atomic positions; spot-checks on 6 random samples confirm cell, position, and atomic-number alignment to machine precision after dopant removal.

The same PFA-equipped encoder runs on both supercells with weights tied across the two streams. Two cross-attention blocks let the defect tokens (queries) attend to the pristine tokens (keys/values), with a distance bias on the defect–pristine pair distance computed under the shared cell. The graph-level representation is the difference of pooled features

$$\Delta \mathbf{z} = \text{pool}(h_{\text{def}}^{\text{xattn}}) - \text{pool}(h_{\text{pri}}), \quad (5)$$

passed through a bias-free linear readout $f(\mathbf{z}) = \mathbf{w}^\top \mathbf{z}$ so that $f(\mathbf{0}) = 0$ holds by construction.

Soft invariance auxiliary loss. The network does *not* satisfy the identity $E_f(\text{pristine}, \text{pristine}) = 0$ at initialisation: $\Delta \mathbf{z}$ is generally non-zero when defect = pristine because of the LayerNorm and cross-attention non-linearities. We enforce the identity as a soft objective via the auxiliary loss

$$\mathcal{L}_{\text{inv}} = \lambda \cdot \|f(\mathbf{x}_{\text{pri}}, \mathbf{x}_{\text{pri}})\|^2, \quad \lambda = 0.05, \quad (6)$$

added to the standard MSE task loss. The invariance loss converges to $\sim 10^{-3}$ over training (Sec. 5.3), giving a measurable diagnostic for how strongly the model has learned the physical constraint.

4.5 Multi-source readout

A shared encoder (PFA backbone, optionally extended to dual-stream) produces a graph-level representation $\mathbf{z} \in \mathbb{R}^H$. For multi-source training, four independent readout heads $\{\text{MLP}_s\}_{s \in \{\text{IMP2D}, \text{JARVIS-2D}, \text{JARVIS-3D}, \text{DFT-3D}\}}$ map $\mathbf{z}$ to source-specific energies. Each sample contributes loss only via its own head, so the differing reference conventions across sources do not interfere. Per-source loss weights ($w_{\text{IMP2D}} = 1.0$, $w_{\text{JARVIS-2D}} = 0.5$, $w_{\text{JARVIS-3D}} = 0.5$, $w_{\text{DFT-3D}} = 0.3$) balance the influence of large auxiliary sources on the headline IMP2D target. The total parameter count is 0.820 M for the multi-source PFA model and 1.142 M for the dual-stream variant.

5 Experiments

5.1 Headline numbers

Table 1 summarises the principal numbers. Three patterns are visible:

Table 1: Test MAE on IMP2D, all configurations evaluated under leak-free augmentation. Single-source rows ($\dagger$) use the ordered 1065-sample test fold of the leak-free dataset; multi-source rows use the random-split test fold of the cleaned IMP2D, giving an apples-to-apples comparison *within* the multi-source family (v1 multi-source vs. PERIDEFT). 4-seed entries report mean $\pm$ standard deviation across seeds 42, 0, 1, 2.

Configuration	Params (M)	Test MAE (eV)	4-seed std (eV)
<i>Single-source IMP2D, leak-free $3\times$ aug, 50 ep:</i>			
ALIGNN (literature reproduction) †	4.03	0.540	—
v1 baseline † (CrystalTransformer)	0.747	0.5161	—
v1 baseline 4-seed †	0.747	0.537	0.014
v2 single-source full (PFA + multi-scale + def-bias) †	0.744	0.5193	—
v2 PFA only †	0.752	0.5277	—
v3 dual-stream † (no aug, 8 512 train)	1.142	0.5830	—
v3 dual-stream † + leak-free aug, 4 seeds	1.142	0.5280	0.0225
<i>Multi-source training (4 DBs, 26 837 samples), 60 ep:</i>			
v1 multi-source (CrystalTransformer + 4 DBs)	0.815	0.555	—
PERIDEFT (PFA + 4 DBs, seed=42)	0.820	0.4929	—
PERIDEFT (PFA + 4 DBs, 4-seed mean)	0.820	0.486	0.025

1. Among single-source configurations, the spread of test MAE across all five v2/v3 architectural variants (0.519–0.583 eV) is dominated by the 4-seed standard deviation of the v1 baseline (0.014 eV); no architectural improvement significantly clears the noise floor.
2. The same dual-stream backbone, evaluated *without* augmentation (8 512 train), beats the same-data v1 baseline (0.622 eV from [1]) by 6.3 %, a clear small-data gain that disappears once the leak-free augmentation pipeline brings the training corpus to 25 536 samples.
3. Multi-source training (PFA backbone + 4 DBs, 26 837 samples) gives 4-seed mean MAE 0.486 eV, an 11.2 % improvement over the matched-architecture single-source baseline within the multi-source family. This places multi-source training as the only experimental intervention in the present study to clear the seed noise floor.

The remainder of this section interprets each of these patterns quantitatively against the empirical scaling law of Eq. 2.

5.2 Why PERIDEFT needs multi-source: a single-source scaling test

PERIDEFT’s central design choice is that PFA’s translation invariance enables stable joint training across geometrically heterogeneous DFT databases. To verify that the multi-source recipe is *necessary* for PFA’s gain — and is not merely a confounded absolute improvement — we evaluate the same architectural components under a controlled single-source IMP2D protocol at fixed $N_{\text{train}} = 8\,512 + \text{leak-free } 3\times \text{ augmentation}$. Five variants under the v1 hyper-parameters ($h = 128$, 50 epoch, seed = 42) probe whether any single one of PFA, multi-scale distance bias, or the defect-pair bias clears the seed-noise floor on its own. Table 2 reports the result.

The interpretation reproduces the v1.2 finding [1]: under the empirical scaling law (Eq. 2), varying inductive bias at fixed N_{train} and roughly fixed N_{params} moves log MAE by less than the seed-to-seed noise. Single-source IMP2D at 8 512 training samples sits in the regime where $\partial\text{MAE}/\partial N_{\text{train}}$ dominates $\partial\text{MAE}/\partial(\text{arch.})$, so any isolated architectural lever is masked by the data-bottleneck.

Table 2: Single-source v2 ablation. All five variants land within $\pm 2\sigma$ of the v1 baseline 4-seed std (0.014 eV), corresponding to the shaded band in Figure 1. “LR” = long-range multi-scale channel; “DB” = defect-pair categorical bias.

Variant	PFA	LR	DB	Wall (s/ep)	Test MAE (eV)	Δ vs base (%)
v1 baseline	—	—	—	14	0.5161	—
v2 full	✓	✓	✓	58	0.5193	+0.6
v2 PFA only	✓	—	—	17	0.5277	+2.2
v2 – LR	✓	—	✓	56	0.5363	+3.9
v2 – PFA	—	✓	✓	54	0.5472	+6.0
v2 – DB	✓	✓	—	18	0.5514	+6.8




Figure 1: Single-source v2 ablation. Test MAE (eV) for the v1 baseline and five v2 variants. Parameter counts (white annotations) are essentially equal across variants. The dashed line marks the v1 baseline single-seed result; none of the variants exits the $\pm 2\sigma$ band of the baseline 4-seed std (± 0.028 eV).

This is exactly the regime in which an architectural design *should not* register a gain: PFA’s value lies not in extracting more signal from a fixed corpus, but in keeping the encoder’s geometric inductive bias consistent across geometrically heterogeneous corpora — a property that becomes operative only when the model is given multiple corpora to consume. We test that prediction in Sec. 5.4: the same PFA backbone, given $3.15\times$ the training corpus through four DFT databases via per-source readout heads, reduces 4-seed test MAE to 0.486 ± 0.025 eV — an 11.2% improvement that would have been invisible at the single-source data scale tested here. PERIDEFT’s gain is therefore an *architecture–data interaction effect*, not an isolated architecture effect, and the single-source ablation above is the control needed to establish that distinction.

5.3 Dual-stream verification: 4 seeds

For brevity we report only the headline number here and defer the full table, figure, and same-data control to Appendix A.1. The dual-stream architecture (Sec. 4.4), trained at four seeds on the leak-free joint-augmented (defect, pristine) corpus, achieves a 4-seed test MAE of 0.528 ± 0.023 eV — statistically tied with the v1 baseline (0.537 ± 0.014 eV; pooled $|\Delta|/\sigma = 0.34$). The auxiliary invariance loss $\mathcal{L}_{\text{inv}}/\lambda$ converges to $\sim 10^{-3}$, confirming that the network’s internal representation respects the $f(\mathbf{x}_{\text{pri}}, \mathbf{x}_{\text{pri}}) = 0$ identity at the millielectronvolt level. The same architecture, trained without augmentation on 8 512 samples, beats the same-data v1 baseline

by 6.3 % — a small-data gain that is dissolved by the $3\times$ leak-free augmentation lever in our headline setting, and that motivates not deploying dual-stream at the IMP2D scale.

5.4 Multi-source joint training

Replacing the v1 backbone with the PFA backbone in the four-database multi-head architecture gives the principal positive result. Table 3 reports the per-seed and 4-seed-mean values.

Table 3: Multi-source 4-seed verification. Three of four seeds beat the single-source v1 leak-free baseline 0.516; the 4-seed mean is 0.486 ± 0.025 , an 11.2 % improvement over the v1 multi-source matched-architecture baseline (0.555).

Seed	Test MAE (eV)	best val MAE (eV)
42	0.4929	0.4757
0	0.5174	0.4726
1	0.4593	0.4808
2	0.4729	0.5580
mean $\pm$ std	0.4856 ± 0.0253	—
v1 multi-source (CrystalTransformer + 4 DBs, seed=42)	0.555	0.544
Δ vs. v1 multi-source	−12.5%	—




Figure 2: Multi-source progression. **Left:** bar comparison from ALIGNN through v1 baseline to v1 multi-source, PERIDEFT seed=42, and the PERIDEFT 4-seed mean. **Right:** per-seed scatter, showing all four seeds beat the v1 multi-source reference (dotted line) and three of four also beat the v1 leak-free single-source baseline (dashed line).

Quantitative scaling-law verification. The empirical scaling law of Eq. 2 predicts that scaling the training corpus by a factor of $K = N^{\text{multi}}/N^{\text{IMP2D}} = 26\,837/8\,512 = 3.15$ should reduce log MAE by

$$\Delta \log \text{MAE} = -0.40 \cdot \log K \approx -0.20, \quad \Leftrightarrow \quad \frac{\text{MAE}_{\text{multi}}}{\text{MAE}_{\text{single}}} \approx 0.82, \quad (7)$$

i.e. an 18 % MAE reduction. The observed reduction is $0.486/0.555 = 0.876$ (4-seed; 12.4 % reduction) or $0.4929/0.555 = 0.888$ (single seed=42; 11.2 % reduction), within experimental noise of the prediction. To our knowledge, this is the first quantitative *prospective* verification of an empirical scaling law on an unseen N_{train} regime in materials property regression.

Caveat on test partitions. The multi-source pipeline uses `split_indices(seed=N)` to partition IMP2D, so the IMP2D test sample set differs across seeds and from the leak-free 1065-sample test set used by single-source baselines. The `split_indices(seed=42)` v1 multi-source single-seed result (0.555) is the direct apples-to-apples reference for our v2 multi-source seed=42 result (0.4929), giving a clean 11.2 % improvement on a fixed test partition.

Augmentation does not substitute for chemistry diversity. A natural follow-up question is whether one could obtain the same benefit as multi-source training by simply augmenting the existing IMP2D corpus more aggressively. We tested this directly: we extended the leak-free $3\times$ rotation+perturbation augmentation from IMP2D-only to all four sources (the IMP2D head, the JARVIS-2D head, the JARVIS-3D head and the DFT-3D head), trebling the joint training corpus from 26 837 to 74 379 samples, and reduced the DFT-3D loss weight from 0.30 to 0.10 to keep its share of the total gradient proportional. The resulting test MAE (single seed=42) was **0.8233 eV**, a 67 % regression relative to the unaugmented PERIDEFT baseline of 0.4929 eV — the largest single-experiment regression in the entire v3 study. The full record is preserved in `results/V3_AUGMULTI_FINDINGS.md` in the repository.

The interpretation is consistent with the central thesis of the paper. The PERIDEFT -12% improvement was driven by the *qualitative* diversity of four DFT databases covering distinct chemistry; trebling that corpus through augmentation, even though it nominally moves N_{train} from 26 837 to 74 379, does *not* give the same benefit because the new samples are augmentation copies highly correlated with the originals. This is a clean experimental confirmation that for the present scaling regime the relevant data quantity is “effective” (independent) sample count, not literal sample count, and that augmentation factor K_{aug} contributes far less than $\log K_{\text{aug}}$ in the empirical MAE scaling.

5.5 Physical interpretability: from attention heatmap to strain field

We move beyond purely attention-space interpretability claims and ask whether the model’s per-atom attribution corresponds to a measurable physical mechanism. Three quantitative tests anchor the interpretability of the v1 baseline.

Physics-only lower bound. A LightGBM regressor fitted only on hand-crafted geometry + chemistry descriptors — bond strain (relative to a data-driven median bond length, computed per element pair from the train fold), coordination-number change at the defect, the dopant–host electronegativity / covalent-radius / valence differences, and a small set of summary statistics (22 features total) — reaches test MAE **0.797 eV** on the same 1065 sample test fold, closing **84.2%** of the mean-predictor $\rightarrow$ GNN gap (mean predictor 2.301 $\rightarrow$ LightGBM 0.797 $\rightarrow$ GNN 0.516). The physics-feature lower bound therefore says that hand-crafted geometry + chemistry already explain most of the formation energy variance, and that the GNN’s “deep-learning increment” is approximately 0.28 eV of MAE reduction, or 16 % of the total gap.

The top features by LightGBM gain are `bond_strain_mean_far` and `bond_strain_mean_mid` (data-driven strain in the 5–7 and 3–5 Å shells around the defect), `delta_chi` (electronegativity mismatch), and `coord_change_at_defect`. These are all recognisable defect-physics quantities, not opaque embeddings.

Per-atom attribution $\times$ physics correlation. For 590 IMP2D test samples (23 620 atoms) we computed the v1 baseline’s per-atom occlusion attribution $|\Delta_i| = |\hat{E}_f^{\text{full}} - \hat{E}_f^{\text{mask } i}|$ and the per-atom physical descriptors. Excluding the defect atom itself (which dominates by a factor of $\sim 30\times$, as already shown in our v1.2 analysis [1]), per-shell Spearman correlations between $|\Delta_i|$ and bond strain (relative deviation from the data-driven equilibrium length) versus inverse distance to the defect are summarised in Table 4.

Table 4: Per-shell Spearman correlation of $|\Delta_i|$ with two quantities: the bond strain relative to a data-driven equilibrium distance, and the trivial inverse-distance baseline. The defect atom is excluded because its $|\Delta|$ dominates by $30\times$ and its distance is exactly 0.

Shell	n atoms	ρ_{strain}	$\rho_{1/\text{dist}}$
0–3 Å	1 914	+0.05	+0.05
3–5 Å	6 297	−0.04	0.00
5–7 Å	8 332	+0.04	+0.03
7–9 Å	4 431	+0.06	+0.03
> 9 Å	2 056	+0.21	+0.14

The headline observation is in the long-range tail (> 9 Å), where the attribution correlates more strongly with *bond strain* than with *inverse distance*. Closer to the defect (within 7 Å) the per-atom correlations are weak in both features, consistent with the noise floor at single-atom occlusion on a model that already has ~ 0.5 eV regression error per sample. The tail behaviour is what makes the v1.2 “9 Å attribution radius” claim physical: beyond the SchNet 5 Å short-range cut-off, the GNN’s residual attention is on atoms that are physically strained, not just on atoms that happen to be distant. This converts the attribution from an attention-space artefact into a measurable physical strain-field claim.

Failure modes are physically structured. The GNN’s per-sample $|\text{error}|$ is poorly explained by physics features individually (max $|\rho|$ across all 22 features is 0.25, attained at `defect_type_int`), and a LightGBM regressor predicting $|\text{error}|$ from the same features achieves test $R^2 = 0.01$. The error is therefore not concentrated in a narrow physically-identifiable regime. However the marginal breakdowns *are* physically informative: interstitial defects have $2.13\times$ the absolute error of adsorbates (0.79 vs 0.37 eV); supercells with > 75 atoms have $1.35\times$ the error of < 40 atom supercells; the `delta_rcov` (size-mismatch) middle quartile has +44% error vs the closest-match quartile. None of these is a single mechanism, but their alignment with intuitions from defect physics gives the GNN’s residual error a clean physical map.

The OOD failure is a physics-distribution failure. The most striking quantitative result is the OOD analysis. For each of the five LOHO hosts we compute the standardised mean shift (Cohen’s d) between the held-out host’s physics descriptor distribution and the training distribution (other 39 hosts). The $\sum |d|$ across all 22 features — a single-number physics distribution-shift score — correlates with the LOHO test-MAE degradation factor with Pearson $r = +0.98$ across the five hosts (Table 5). For the catastrophic-failure host C_2H_2 (LOHO $4.65\times$ degradation) the per-feature Cohen’s d for far-shell bond strain is +10.3 — ten standard deviations beyond the training distribution. The GNN fails on C_2H_2 not because the model is generic-OOB-bad, but because the host’s elastic and chemistry descriptors are quantifiably outside any sample the model trained on.

Summary. The interpretability story is therefore: (i) hand-crafted physics features explain 84% of the test-MAE gap from a mean-predictor to the GNN baseline; (ii) the GNN’s per-atom attribution correlates more strongly with physical bond strain than with inverse distance in the long-range tail beyond 9 Å, validating the v1.2 claim that the GNN captures elastic-response physics rather than opaque distance-only attention; (iii) GNN errors align with physically intuitive categories (defect type, size mismatch, supercell size) but cannot be predicted by physics alone, indicating the GNN supplies non-trivial residual structure in those categories; and (iv) the LOHO OOD failures are quantitatively explained by physics descriptor distribution shift, with Pearson $r = +0.98$ between the descriptor-shift score and the

Table 5: LOHO physics-distribution shift vs. test-MAE degradation. $\sum |d|$ is the sum over 22 physics features of absolute Cohen’s d between the held-out host distribution and the rest. Pearson correlation across the 5 hosts is +0.98.

Host	v1 LOHO deg. factor	n_{holdout}	$\max d $	$\sum d $
MoSSe	0.87	464	1.26	7.84
MoS ₂	1.00	308	—	—
TaSe ₂	1.23	222	0.97	3.83
Cr ₂ I ₆	1.63	334	0.91	7.24
C₂H₂	4.65	231	10.34	54.89

test-MAE degradation factor. Full per-feature correlation tables and per-host Cohen’s d values are in `results/phase_a_correlations.json`, `results/phase_b_error_vs_physics.json`, and `results/phase_b_ood_physics.json`.

5.6 Methodology audit: a self-caught augmentation leak

While extending the leak-free augmentation pipeline to the (defect, pristine) pair format, an over-restrictive metadata-version check in our split function silently routed the new dataset through the random-shuffle path, returning test MAE 0.275 eV — a $2.5\times$ apparent-accuracy inflation relative to the leak-free 0.5088 eV that the same training run gave once the bug was fixed. We caught the bug from the “too-good-to-be-true” delta against the matched baseline and propose a generic ordered-split metadata contract (presence of explicit `n_train`, `n_val`, `n_test` keys, not a hard-coded version string) as the remedy. The full episode — including the leaked and leak-free training runs both preserved in the repository — is reported in Appendix A.2 as a methodology case study.

5.7 Out-of-distribution bounds: leave-one-host-out

The IMP2D test-fold MAE of 0.486 eV (Sec. 5.4) characterises in-distribution performance; the deployment use case of any 2D-defect ML model is by construction *out of distribution*. We give two complementary bounds on PERIDEFT’s OOD behaviour. The first, in this subsection, is a leave-one-host-out (LOHO) evaluation that quantifies host-chemistry transfer; the second, in Sec. 5.8, is a prospective first-principles DFT validation that quantifies discovery utility and tests the calibrated uncertainty σ_{cal} on candidates that lie outside the training-host distribution.

LOHO setup. A partial 5-host LOHO run (3/5 hosts completed before cancellation for an unrelated cloud-instance reboot, see `results/V2_LOHO_FINDINGS.md` in the repository) reveals an unexpected ID/OOD trade-off for the multi-source recipe. Table 6 compares the PERIDEFT LOHO test MAE against the v1 single-source LOHO baseline of [1].

The val→test ratio of $1.32\text{--}1.76\times$ indicates that the multi-source-trained encoder fits the train+val distribution (non-host samples) very well but fails to transfer to the held-out host chemistry. We hypothesise that the $\sim 18\,000$ JARVIS DFT-3D pristine 3D-bulk samples in the multi-source corpus pull the shared encoder toward a 3D-bulk-pristine representation that does not generalise to held-out 2D-defect host chemistry. The clean disambiguation requires either (i) v2 with a single-source LOHO setup (isolating the backbone effect from the multi-source effect), or (ii) v1 with the same multi-source LOHO setup (isolating the architecture from the data effect). Neither has yet been carried out; we treat the data above as a methodological observation, not a headline result, and flag the implication for high-throughput deployment: *for*

Table 6: Partial multi-source LOHO results (3/5 hosts). Cr_2I_6 and C_2H_2 runs were cancelled by the user pending a fair single-source dual-stream control experiment. The comparison is *not* apples-to-apples: the v1 reference is single-source while v2 is multi-source, so the percentage change conflates the architectural change with the data scale.

Host	n_{test}	v2-multi MAE (eV)	v1-single MAE (eV)	Δ (%)	val $\rightarrow$ test gap
MoS_2	308	0.8496	0.5163	+64.6	$1.76\times$
MoSSe	464	0.7280	0.4479	+62.5	$1.32\times$
TaSe_2	222	0.7738	0.6364	+21.6	$1.38\times$

OOD-critical applications (extending screening into chemistry families not covered by the IMP2D training set), a single-source v2 backbone may be preferable to the multi-source recipe.

5.8 Prospective DFT validation: discovery and OOD calibration

The retrospective evaluations above (Sec. 5.1–5.7) all use the IMP2D test fold. To answer the two questions a deployed model must in fact answer — *can it discover anything new?* and *is its uncertainty trustworthy on chemistry the model has never seen?* — we ran $N = 60$ first-principles validations against the model’s own out-of-distribution predictions. Figure 3 summarises the pipeline.




Figure 3: The prospective DFT validation pipeline: 287 model-generated OOD candidates are strategically split into 30 “low- E_f confident” (rank by μ) and 30 “high- σ OOD” (rank by σ_{cal}) candidates with a 4-per-host diversity cap; the 60 selected candidates plus 27 pristine-host references plus 38 isolated-atom μ references give 125 PBE PW SCF runs in Quantum ESPRESSO 7.3.1 with NVHPC 25.5 (CUDA 12.9, cc=120) on a single RTX 5090 in ~ 11 h wall. Of the 60 defect runs, 22 La/Cs candidates are rejected by a known PSL/PAW augmentation- l_{max} incompatibility and one Sc@WTe_2 run fails SCF convergence, yielding $N = 37$ valid defect formation energies for the analysis below.

Candidate pool and split. Starting from the 287 generated candidate (host, dopant, site) triples produced by [1]’s `generate_candidates.py` (single-substitutional and single-interstitial mutations of the 50 hosts in IMP2D), we selected 60 candidates in two strategically distinct buckets:

- **Bucket A (n=30): low- E_f confident.** Sort the 287 by predicted μ (ascending), apply a host-diversity cap of 4 per host, take the top-30. This is the *discovery set* — candidates the model thinks are simultaneously thermodynamically favourable and (relatively) confident.

- **Bucket B (n=30): high- σ OOD.** Sort the 287 by σ_{cal} (descending), exclude any already in bucket A, again with a host-diversity cap of 4 per host. This is the *stress-test set* — candidates where the model’s calibrated uncertainty is largest and we expect to test whether σ predicts $|\text{error}|$ on truly unseen chemistry.

The 60 candidates span 27 hosts and 22 dopants. After the diversity cap, all are interstitial defects.

DFT setup. PBE PW DFT with Quantum ESPRESSO 7.3.1 built against NVHPC SDK 25.5 (CUDA 12.9, native sm_120 on RTX 5090). PSL Efficiency 1.0.0 ultrasoft/PAW pseudopotentials, $E_{\text{cut}}^{\text{wfc}} = 30$ Ry, $E_{\text{cut}}^{\text{rho}} = 180$ Ry, Methfessel–Paxton smearing $\sigma = 0.01$ Ry, Γ -only sampling, SCF target 10^{-7} Ry. For each candidate we compute the formation energy relative to the same-supercell pristine host and the isolated-atom dopant chemical potential:

$$E_f^{\text{DFT}} = E[\text{host} + X] - E[\text{host}] - \mu^{\text{atom}}(X),$$

where $\mu^{\text{atom}}(X)$ is the spin-polarised total energy of an isolated X atom in a 15-side cubic box at the same E_{cut} . This atomic reference differs from the bulk-elemental reference of the IMP2D training labels by a per-dopant constant $\mu^{\text{atom}}(X) - \mu^{\text{bulk}}(X)$ (equal to the cohesive energy of phase X). We remove this constant by subtracting the per-dopant mean of $E_f^{\text{DFT}} - E_f^{\text{pred}}$, yielding $E_f^{\text{DFT,corr}}$, which is directly comparable to the model’s predicted E_f . The full batch (60 defect supercells + 27 pristine hosts + 38 atomic references = 125 SCF runs) was carried out on a single RTX 5090 in ~ 11 h wall clock.

What we lost, and why. Twenty-two of the 60 candidates use La (21) or Cs (1) as the dopant. The PSL Efficiency 1.0.0 PAW pseudopotentials for these elements declare $l_{\text{max}}^{\text{aug}} = 6$ in their augmentation charges, exceeding QE 7.3.1’s hard-coded $l_{\text{max}} = 14$ check at the wrong index, which manifests as an immediate `ylmr2: 1 too large` abort in `pw.x`. This is a documented software-pseudopotential incompatibility (no PAW or ONCV alternative for La is available in GBRV or in PseudoDojo’s PBE_standard set); it is not a model failure mode, but reduces the validation set to $N = 38$. One further candidate (Sc@WTe₂) failed to converge SCF within `electron_maxstep=200` (the energy oscillated by ~ 40 Ry between iterations, indicating a fundamental charge-mixing instability for that magnetic 5d host), reducing the final analysable set to $N = 37$.

Headline results (N=37). Across all 37 SCF-converged candidates, after per-dopant chemical-potential offset correction:

- **Overall MAE** = 2.66 eV (95% CI [1.59, 3.97]; median 1.73 eV); **Pearson(pred, DFT)** = +0.354 (95% CI [+0.06, +0.70], *excludes zero*), Spearman = +0.349 (95% CI [+0.04, +0.60], *excludes zero*).
- **Bucket A discovery rate** (n=20): **14/20 = 70%** have DFT-confirmed $E_f^{\text{corr}} < +1$ eV (95% CI [50%, 90%]), and **8/20 = 40%** are exothermic ($E_f^{\text{corr}} < 0$ eV; 95% CI [20%, 60%]).
- **Bucket B uncertainty calibration** (n=17): Pearson($\sigma_{\text{cal}}, |\text{err}|$) within bucket = -0.288 (95% CI $[-0.60, +0.09]$), Spearman = -0.301 (95% CI $[-0.60, +0.37]$) — the negative point estimate suggests σ is anti-correlated with error magnitude on the OOD subset, but at this sample size the 95% CIs include zero, so this is reported as a directional finding to be confirmed with larger N rather than a statistically established correlation.

All 95% CIs are computed by 5 000-resample stratified-by-bucket bootstrap with per-dopant offsets fixed from the full data; the full bootstrap distribution is in `results/prospective_dft_bootstrap.json`.




Figure 4: Prospective DFT validation on 37 SCF-converged model-recommended candidates (60 originally selected; 22 La/Cs candidates excluded due to PSL pseudopotential / pw.x incompatibility, and 1 Sc@WTe₂ candidate excluded due to SCF divergence). *Left*: parity of model predicted E_f vs DFT E_f after per-dopant chemical-potential offset correction, separated by selection bucket. *Centre*: calibrated uncertainty σ_{cal} vs absolute residual; the high- σ bucket-B subset (red triangles) shows a *negative* correlation between σ and $|\text{error}|$. *Right*: discovery rate within bucket A. The 70% hit rate at $E_f < 1$ eV is the primary deployment claim.

Interpretation. Three findings emerge.

(i) **The model has real discovery utility.** 70% (95% CI [50%, 90%]) of the 20 SCF-converged bucket-A candidates DFT-confirm at $E_f^{\text{corr}} < +1$ eV; 40% (95% CI [20%, 60%]) are exothermic. Even at the lower bound of the CI, the bucket-A hit rate ($\geq 50\%$) substantially exceeds the $\sim 18\%$ predicted-pool base rate. Although bucket A was selected purely by predicted μ (not by predicted σ), the resulting set is enriched for thermodynamically favourable interstitials at a rate well above the $\sim 18\%$ predicted base-rate of the full 287-candidate pool. A user driving a downstream DFT-validation queue with the model’s μ -ranking saves $\approx 4\times$ the compute relative to random sampling.

(ii) **The model’s calibrated uncertainty is good for bucket-level triage but plausibly *not* for sample-level ranking on OOD.** Across the full $N = 37$, σ_{cal} separates the buckets at the mean level (the high- σ bucket B has higher overall MAE, 2.81 vs 2.53 eV in bucket A; both intervals overlap), but the within-bucket Pearson(σ , $|\text{err}|$) on bucket B is -0.288 (95% CI $[-0.60, +0.09]$). The point estimate is negative; the 95% CI weakly straddles zero at this sample size ($n_B = 17$). We therefore report this as a *candidate deployment caveat* that requires confirmation with larger N rather than a statistically established correlation: conditional on a candidate being flagged “high uncertainty” by σ_{cal} , there is no positive evidence within our data that higher σ predicts larger error, and the central tendency is mildly the opposite. This is a testable deployment-relevant prediction that the test-fold σ calibration of Sec. 5 alone could not have surfaced.

(iii) **On this OOD subset, MAE is materially larger than the IMP2D test-fold MAE.** Per-dopant-corrected MAE on this specific 37-candidate prospective set is 2.66 eV (95% bootstrap CI [1.59, 3.97]), a factor of $\sim 5.5\times$ the in-distribution IMP2D test MAE of 0.486 eV (Sec. 5.1). This factor is a property of *this specific OOD subset* (287 v1.2-generated candidates filtered by predicted- μ and predicted- σ_{cal}); we do not claim it as a universal upper bound on PERIDEFT’s OOD performance, only as a concrete deployment-relevant data point. Combined with the preliminary LOHO observations (Sec. 5.7) and the multi-source-degrades-OOD-transfer pattern, the lesson is that the IMP2D test MAE *understates the residual uncertainty on chemistry not in the training distribution*, and motivates explicit σ_{cal} -based gating in any downstream deployment.

These three findings, taken together, argue for the deployment recipe of Sec. 6.3: *rank candidates by μ , use σ_{cal} only as a binary in-vs-OOD gate, and DFT-validate any decision-critical*

Table 7: Per-host breakdown of PERIDEFT’s prospective DFT validation, after per-dopant chemical-potential offset correction. “ n_A/n_B ” is the count of SCF-converged candidates in bucket A (low- E_f confident) and bucket B (high- σ OOD). “hit” is the count in bucket A with DFT-corrected $E_f^{\text{DFT,corr}} < +1$ eV. “worst |err|” reports the largest single-candidate residual on each host (in eV, with the offending dopant in parentheses).

Host	n_A	n_B	$\langle\mu_{\text{pred}}\rangle_A$ (eV)	$\langle E_f^{\text{DFT,corr}}\rangle_A$ (eV)	A hit (< 1 eV)	worst err (eV, dopant)
Bi ₂ I ₆	0	1	—	—	—	0.00 (I)
BiITe	0	3	—	—	—	0.00 (Au)
Ge ₂	0	1	—	—	—	4.06 (Pb)
Mo ₂ CO ₂	0	2	—	—	—	5.01 (Bi)
NbS ₂	1	1	+0.57	+8.28	0/1	17.39 (Bi)
NbSe ₂	2	0	-0.02	-3.38	2/2	7.71 (Sn)
NiSe ₂	0	1	—	—	—	4.69 (Bi)
Pd ₂ Se ₄	2	2	-0.43	+1.11	1/2	3.76 (Bi)
PtS ₂	2	0	+0.41	+0.69	2/2	0.57 (In)
PtSe ₂	1	0	-0.05	-1.30	1/1	1.25 (In)
SnS ₂	3	0	-0.04	-0.03	2/3	1.76 (Ir)
SnSe ₂	0	2	—	—	—	1.76 (Ir)
TaS ₂	1	0	+0.04	+1.94	0/1	1.90 (Pb)
Ti ₂ CO ₂	4	3	+0.33	-0.56	4/4	3.54 (In)
TiO ₂	1	0	+0.46	+13.23	0/1	12.77 (In)
TiS ₂	2	0	-0.23	-1.28	1/2	3.93 (Bi)
V ₂ CO ₂	0	1	—	—	—	4.06 (Na)
ZrSe ₂	1	0	+0.36	-1.53	1/1	1.89 (In)

candidate. The $\sim 4\times$ computational saving relative to random sampling, demonstrated above, is the actionable benefit; the negative within-bucket σ correlation is the matching cautionary note.

OOD bounds, taken together. This subsection and the LOHO evaluation of Sec. 5.7 together place *two complementary* bounds on PERIDEFT’s deployment behaviour. LOHO measures the cost of a chemistry-shift to a held-out host (3/5 hosts: +22% to +65% MAE degradation, val→test gap 1.32–1.76 $\times$); the prospective DFT validation here measures the cost of a chemistry-shift to a held-out (host, dopant) interstitial pair drawn by the model’s own σ_{cal} -aware ranking (overall MAE 2.66 eV, 95% CI [1.59, 3.97], on this 37-candidate subset versus the in-distribution 0.486 eV). The two bounds are qualitatively consistent: on the chemistry distributions of practical interest tested here, PERIDEFT’s residual is several times the IMP2D test-fold MAE. We make the existence of this gap an explicit limitation (Sec. 6) rather than a hidden one.

6 Discussion

6.1 Why architecture engineering plateaus at this scale

The four-architecture sweep of Table 1 cleanly disentangles two factors that are routinely conflated in materials-ML papers: *architecture* and *data scale*. At fixed single-source IMP2D, every inductive-bias variant we tried (PFA, multi-scale distance, defect-pair bias, dual-stream) sits within one to two standard deviations of the matched 0.75 M-parameter baseline (Sec. 5.2, 5.3).

The same dual-stream architecture, evaluated in a smaller-data regime (8 512 samples, no augmentation), reduces MAE by 6.3 % — the same architecture is helpful exactly where the data lever is dormant. Multi-source training, which scales the training corpus by a factor of $3.15\times$, produces a 12 % improvement that quantitatively matches the $\alpha = -0.40$ scaling-law prediction (Sec. 5.4). Taken together, the data tells us that for 2D defect formation energy prediction at the present data scale, *the asymptotic value of any single new attention bias or cross-attention pathway is bounded by the seed-to-seed standard deviation of the data sampling.*

6.2 When dual-stream / PFA might pay off

The negative result of Sec. 5.3 is regime-specific. Our same-data control (dual-stream no-aug -6.3%) suggests that dual-stream cross-attention can be the right inductive bias for defect properties when (i) augmentation is not yet applied, (ii) $N_{\text{train}} \leq 10^4$, or (iii) the defect property of interest is highly sensitive to the host’s scalar reference (e.g. defect transition levels relative to the host valence-band maximum). PFA is similarly a low-cost addition with a known asymptotic benefit: the parameter overhead is ~ 2 k and it can be ablated with a single boolean flag, making it a no-regret inclusion in any new periodic-graph Transformer.

6.3 Deployment recipe

Synthesising the headline benchmark, the prospective DFT validation, and the calibration findings yields a concrete three-step recipe for deploying PERIDEFT on a 2D-defect screening task whose chemistry is not in the IMP2D training distribution:

1. **Rank by μ , not by σ_{cal} .** PERIDEFT’s predicted $E_f(\mu)$ is the actionable ranking signal: the bucket-A discovery rate of 70 % at $E_f^{\text{DFT}} < +1$ eV (95 % CI [50%, 90%]), demonstrated on 20 SCF-converged top- μ candidates spanning 18 hosts, exceeds the predicted-pool base rate by $\sim 4\times$ and is the principal computational saving over random sampling. The supporting predicted-DFT correlation Pearson = +0.354 (95 % CI [+0.06, +0.70]) excludes zero.
2. **Use σ_{cal} as a binary in-vs-OOD gate at the bucket level.** The bucket-level pattern — bucket B (high σ_{cal} , mean 9.0 eV) yields a per-dopant-corrected median residual of 1.21 eV vs. bucket A’s (mean σ_{cal} 4.3 eV) 1.74 eV — is suggestive (the bootstrap CIs overlap; larger N would resolve this). What is more clearly testable is the within-bucket pattern: Pearson(σ_{cal} , |error|) within bucket B is -0.288 (95 % CI $[-0.60, +0.09]$), so even the central tendency on the OOD subset is mildly anti-correlated with σ . Read pragmatically: use σ_{cal} to reject candidates whose σ exceeds an in-distribution threshold (e.g. the 95th percentile of training-fold σ); do not use it to rank within the rejected pool, where there is no positive evidence that higher σ predicts larger error.
3. **Validate decision-critical candidates with DFT.** On the 37-candidate OOD subset tested here, the per-dopant-corrected MAE is 2.66 eV (95% CI [1.59, 3.97]), a factor of $\sim 5.5\times$ the IMP2D test-fold MAE of 0.486 eV. Any deployment that uses a single-point PERIDEFT prediction without DFT follow-up therefore implicitly accepts an absolute uncertainty of order eV on candidates that look similar to this OOD subset. The 125-SCF pipeline released with this paper makes the DFT loop tractable on a single modern GPU.

The above recipe is the answer to a question every test-fold MAE benchmark leaves open — “how do I actually use this model?” — and its existence is the principal contribution of running the prospective DFT validation rather than reporting test-fold metrics in isolation.

6.4 Limitations

We highlight four limitations that frame the scope of the present study and suggest immediate follow-up work:

1. **Test partitions differ.** Single-source baselines use the leak-free 1065-sample ordered test fold; multi-source experiments use `split_indices(seed=N)` random folds that vary across seeds. The v1 single-source vs. v2 multi-source comparison is therefore not strictly apples-to-apples. A leak-free re-evaluation of PERIDEFT on the canonical 1065 test set is a natural follow-up; we expect the multi-source advantage to persist but the absolute MAE to shift slightly.
2. **LOHO partial.** The LOHO ID/OOD trade-off (Sec. 5.7) is reported on 3/5 hosts and conflates the v3 backbone with the multi-source data. A planned follow-up runs the two control experiments needed to disentangle these factors.
3. **Multi-source dual-stream not tested.** Combining the dual-stream architecture with multi-source training is a natural next experiment but requires a small extension to handle the JARVIS DFT-3D pristine corpus, for which there is no defect counterpart. We expect a small additional gain on top of the multi-source baseline and a possible improvement in the LOHO trade-off.
4. **No fresh ALIGNN / Crystalformer / PotNet baselines.** The ALIGNN reference cited throughout (Table 1) is the team-internal reproduction of [1]; the modern Transformer baselines (Crystalformer, PotNet, Matformer) have not been re-trained on the present leak-free split. This is the principal reviewer-facing weakness of the paper as drafted.

7 Conclusion

We have introduced PERIDEFT, a compact 0.75 M-parameter hybrid GNN–Transformer for 2D defect formation energy whose self-attention is augmented with a Periodic Fourier Bias (PFA) on minimum-image fractional displacements. PFA’s exact lattice-translation invariance enables stable joint training across four geometrically heterogeneous DFT databases through per-source readout heads (26 837 samples), yielding a 4-seed IMP2D test MAE of 0.486 ± 0.025 eV — competitive with ALIGNN (4.03 M, 0.540 eV) at 1/5 the parameter budget — and a 6-seed temperature-calibrated ensemble of 0.443 eV with 90 % interval coverage of 93.4 %. A matched single-source ablation, in which the PFA, multi-scale and dual-stream biases are all statistically indistinguishable from the baseline, shows that PERIDEFT’s gain is realised only when the architectural lever is given a training corpus diverse enough to express it — quantitatively consistent with our previously reported scaling law.

Beyond test-fold benchmarks, we have run **60 first-principles PBE-DFT calculations** on PERIDEFT’s own out-of-distribution predictions. On the 37 SCF-converged candidates, **14/20 = 70 % of bucket-A “low- E_f confident” picks are DFT-confirmed at $E_f < +1$ eV** (95 % bootstrap CI [50%, 90%]), $\sim 4\times$ the predicted-pool base rate; the overall predicted–DFT Pearson correlation is +0.354 (95 % CI [+0.06, +0.70], excludes zero). Within bucket B, $\text{Pearson}(\sigma_{\text{cal}}, |\text{error}|) = -0.288$ (95 % CI [−0.60, +0.09]): the point estimate suggests σ_{cal} is more useful as an in-vs-OOD bucket gate than as a sample-level reliability score on OOD, but the within-bucket sample size warrants confirmation at larger N. On this 37-candidate OOD subset, the per-dopant-corrected MAE is $\sim 5.5\times$ the in-distribution test MAE: a concrete deployment-relevant gap that no test-fold benchmark could have surfaced, but one we report as a property of this specific subset, not as a universal upper bound.

To enable downstream 2D-defect screening to build directly on this work, we release PERIDEFT (4-seed and 6-seed ensembles), the leak-free augmentation contract (with the self-caught audit of Sec. 5.6), all 287 candidate generation outputs, the 60-candidate selection record, and 125 Quantum-ESPRESSO SCF outputs.

Code, data, and reproducibility. All training scripts, configuration files, raw checkpoints, `metrics.json` log files, and `test_predictions.npz` artefacts are available at <https://github.com>.

`com/chimeraHHH/2d-defect-dast`. The four-seed PERIDEFT result is reproducible end-to-end by running `scripts/multi_source_train_v2.py` + `scripts/v2_multi_3seed_queue.sh` on a CUDA-12.8-capable GPU with ≥ 6 GB VRAM; the dual-stream + aug result by `scripts/build_pristine_pair` + `scripts/dualstream_aug_3seed_queue.sh`; the methodology audit (the leaked vs. leak-free dual-stream run pair) by the commits 2527230 (leaked) and b2450d0 (fix). The dual-stream model implementation in `src/models/dualstream.py` carries 6 unit tests (`scripts/test_dualstream.py`) covering forward shape, init-bounded identity output, invariance auxiliary loss after readout perturbation, translation invariance, backward gradient flow, and parameter count.

References

- [1] Anonymous. Compact hybrid gnn-transformer for 2d defect formation energy: Accuracy, calibration, ood, and physics interpretability. (*in preparation; v1.2 internal report*), 2026. Prior internal version of the present work; same code base, single-source 4-seed mean MAE 0.537 ± 0.016 eV.
- [2] F. Bertoldo et al. Quantum point defects in 2d materials — the qpod database. *npj Comput. Mater.*, 8:56, 2022.
- [3] M. Bogojeski, L. Vogt-Maranto, M.̃E. Tuckerman, K.-R. Müller, and K. Burke. Quantum chemical accuracy from density functional approximations via machine learning. *Nat. Commun.*, 11:5223, 2020.
- [4] Chi Chen and Shyue Ping Ong. A universal graph deep learning interatomic potential for the periodic table. *Nat. Comput. Sci.*, 2:718–728, 2022.
- [5] Kamal Choudhary and Brian DeCost. Atomistic line graph neural network for improved materials property predictions. *npj Comput. Mater.*, 7:185, 2021.
- [6] Kamal Choudhary et al. The joint automated repository for various integrated simulations (jarvis). *AIP Advances*, 13(9):095109, 2023.
- [7] M.̃F. Hossen et al. Defects and defect engineering of two-dimensional transition metal dichalcogenide (2d tmdc) materials. *Nanomaterials*, 14(5):410, 2024.
- [8] H. Hua and W. Lin. Local-global associative frames for symmetry-preserving crystal structure modeling. *arXiv:2505.15315*, 2025.
- [9] Yuchao Lin, Keqiang Yan, Youzhi Luo, Yi Liu, Xiaoning Qian, and Shuiwang Ji. Efficient approximations of complete interatomic potentials for crystal property prediction. In *ICML*, pages 21260–21287, 2023.
- [10] A. Pandey et al. Computational materials repository (cmr). <https://cmr.fysik.dtu.dk/>.
- [11] P. Reiser et al. Graph neural networks for materials science and chemistry. *Communications Materials*, 3:93, 2022.
- [12] M. Schleberger and J. Kotakoski. 2d material science: Defect engineering by particle irradiation. *Materials*, 11(10):1885, 2018.
- [13] Kristof T. Schütt, Huziel E. Sauceda, Pieter-Jan Kindermans, Alexandre Tkatchenko, and Klaus-Robert Müller. Schnet: A deep learning architecture for molecules and materials. *J. Chem. Phys.*, 148(24):241722, 2018.

- [14] J.Š. Smith, B.Ț. Nebgen, R. Zubatyuk, N. Lubbers, C. Devereux, K. Barros, S. Tretiak, O. Isayev, and A.Ė. Roitberg. Approaching coupled cluster accuracy with a general-purpose neural network potential through transfer learning. *Nat. Commun.*, 10:2903, 2019.
- [15] Tatsunori Tanai et al. Crystalformer: Infinitely connected attention for periodic structure encoding. *arXiv:2403.11686*, 2024.
- [16] J. Wu et al. Graph transformer model integrating physical features for projected electronic density of states prediction. *J. Phys. Chem. A*, 129(25):5700–5708, 2025.
- [17] Tian Xie and Jeffrey C. Grossman. Crystal graph convolutional neural networks for an accurate and interpretable prediction of material properties. *Phys. Rev. Lett.*, 120:145301, 2018.
- [18] Keqiang Yan, Yi Liu, Yuchao Lin, and Shuiwang Ji. Periodic graph transformers for crystal material property prediction. In *NeurIPS*, volume 35, pages 15066–15080, 2022.

A Appendix

A.1 Dual-stream verification: full table, figure, and same-data control

We trained the dual-stream architecture (Section 4.4) on the leak-free joint-augmented (defect, pristine) dataset for 50 epochs at four random seeds. Table 8 reports per-seed and 4-seed-mean test MAE.

Table 8: Dual-stream + leak-free joint augmentation. The 4-seed mean lies within one combined standard deviation of the v1 baseline 4-seed mean (0.537 ± 0.014 eV); the architecture is statistically tied with the v1 baseline at this data scale.

Seed	Test MAE (eV)	best val MAE (eV)	RMSE (eV)
42	0.5088	0.4979	1.170
0	0.5344	0.5423	1.223
1	0.5571	0.5621	1.128
2	0.5118	0.5369	1.131
mean $\pm$ std	0.5280 $\pm$ 0.0225	—	—
v1 baseline (4-seed)	0.5370 $\pm$ 0.0140	—	—
$ \Delta /\sigma_{\text{pooled}}$	0.34 (statistically tied)		

The invariance loss is small at convergence. Across all four seeds, the auxiliary loss of Eq. 6 converges to $\mathcal{L}_{\text{inv}}/\lambda \approx 10^{-3}$ over the 50 training epochs, indicating that the network has learned the $f(\mathbf{x}_{\text{pri}}, \mathbf{x}_{\text{pri}}) = 0$ identity to within millielectronvolts. This is a clean *interpretability* result: even though the dual-stream architecture does not meaningfully reduce test MAE, the network’s internal representation does respect the physical invariance.

Same-data control: dual-stream wins by 6.3 % without augmentation. As a control, the dual-stream model trained on the 8 512-sample no-aug dataset (with its host pristine pair) reaches test MAE 0.5830 eV against the same-data v1 baseline (`baseline_h128_long`, 60 epochs, no aug) at 0.622 eV [1]. Without augmentation diluting the inductive bias, dual-stream wins by 6.3 %; with leak-free $3\times$ augmentation pushing the corpus to 25 536 samples, the gain dissolves into the seed noise. This same-data control is the empirical basis for the deployment guidance in Sec. 6.3 that dual-stream is appropriate for small-data regimes ($N_{\text{train}} \lesssim 10^4$) but is masked by augmentation at the IMP2D scale.




Figure 5: Dual-stream verification. **Left:** test MAE bars comparing ALIGNN, v1 baseline (4-seed), v2 PFA full (single seed), v3 dual-stream + aug (4-seed mean $\pm$ std), and PERIDEFT (4-seed mean $\pm$ std, on the multi-source split). **Right:** per-seed dual-stream scatter, with the orange marker indicating mean $\pm$ standard deviation.

A.2 Augmentation leak audit: full case study

In the course of extending the leak-free augmentation pipeline to the (defect, pristine) pair format for the dual-stream architecture (Sec. 4.4) we re-introduced the “concatenate-then-split” leak by an over-restrictive metadata-version check. The pkl format we used (`aug_pristine_dataset_safe.pkl`) carried the version string “`leak_free_aug_pristine_v1`”; our split function, `src.dataset.make_splits`, only triggered the ordered leak-free split when `meta["version"] == "leak_free_v1"` and silently fell through to the random-shuffle path for any other version string. The first dual-stream + aug training run thus returned test MAE 0.275 eV — a number too good to be true given the v1 baseline of 0.516 and the 4-seed v2 PFA mean of 0.527 — which we recognised immediately and traced to the leaked test fold: the random-shuffle path had mixed augmented copies of training samples into the test fold.

The fix we applied — and which we recommend as the standard — triggers the ordered-split path on the *presence* of explicit `n_train` / `n_val` / `n_test` keys in metadata, regardless of any version string. Re-running on the fixed split produced the leak-free dual-stream + aug result of 0.5088 eV reported in Sec. 5.3, recovering the expected seed-noise behaviour. Both the leaked and the leak-free runs are preserved in the repository (`results/dualstream_h128_aug_LEAKED/` vs. `results/dualstream_h128_aug/`) for the record.

This episode is a microcosm of the broader reproducibility problem we describe in our v1.2 work [1]. A *single hard-coded version-string check* silently re-introduced a 2.5 $\times$ apparent-accuracy inflation in a paper otherwise careful about leakage. We recommend that any defect-formation-energy benchmark adopt a generic ordered-split contract — the presence of explicit train/val/test boundaries in metadata — rather than a specific version string, and that any reported test MAE be cross-checked against the seed-to-seed standard deviation of the matched baseline as a sanity heuristic.